

```
1. # Useful starting lines
2. %matplotlib inline
3. import numpy as np
4. import matplotlib.pyplot as plt
5. %reload_ext autoreload
6. %autoreload 2
```

```
1. # Check the Python version
2. import sys
3. if sys.version.startswith("3."):
4.     print("You are running Python 3. Good job :)")
5. else:
6.     print("This notebook requires Python 3.\nIf you are using Google Colab, go to
    Runtime > Change runtime type and choose Python 3.")
```

```
1. You are running Python 3. Good job :)
```

```
1. # Load the data
```

```
1. import datetime
2. from helpers import *
3.
4. height, weight, gender = load_data(sub_sample=False, add_outlier=False)
5. x, mean_x, std_x = standardize(height)
6. b, A = build_model_data(x, weight)
```

```
1. print('Number of samples n = ', b.shape[0])
2. print('Dimension of each sample d = ', A.shape[1])
```

```
1. Number of samples n = 10000
2. Dimension of each sample d = 2
```

Least Squares Estimation

Least squares estimation is one of the fundamental machine learning algorithms. Given an $n \times d$ matrix A and a $n \times 1$ vector b , the goal is to find a vector $x \in \mathbb{R}^d$

which minimizes the objective function $f(x) = \frac{1}{2n} \sum_{i=1}^n (a_i^{\top} x - b_i)^2 = \frac{1}{2n} \|Ax - b\|^2$

In this exercise, we will try to fit x using Least Squares Estimation.

One can see the function is L smooth with $L = \frac{1}{n} \|A^{\top} A\| = \frac{1}{n} \|A\|^2$.

Computing the Objective Function

Fill in the `calculate_objective` function below:

```
1. def calculate_objective(Axmb):
2.     """Calculate the mean squared error for vector Axmb = Ax - b."""
3.     # *****
4.     # INSERT YOUR CODE HERE
5.     # TODO: compute mean squared error
6.     # *****
7.     squared = Axmb ** 2
8.     sum_squared = np.sum(squared)
9.     mse = 0.5 * sum_squared / len(Axmb)
10.    return mse
```

Compute smoothness constant L

To compute the spectral norm of A you can use `np.linalg.norm(A, 2)`

```
1. def calculate_L(b, A):
2.     """Calculate the smoothness constant for f"""
3.     # *****
4.     # INSERT YOUR CODE HERE
5.     # TODO: compute ||A.T*A||
6.     # *****
7.     ATA = A.T @ A
8.     # *****
9.     # INSERT YOUR CODE HERE
10.    # TODO: compute L = smoothness constant of f
11.    # *****
12.    L = np.linalg.norm(ATA, 2) / len(Axmb)
13.    return L
```

Gradient Descent

Please fill in the functions `compute_gradient` below:

```
1. def compute_gradient(b, A, x):
2.     """Compute the gradient."""
3.     # *****
4.     # INSERT YOUR CODE HERE
5.     # TODO: compute gradient and objective
6.     # *****
7.     Axmb = A @ x - b
8.     grad = A.T @ Axmb / len(Axmb)
9.     return grad, Axmb
```

Please fill in the functions `gradient_descent` below:

```
1. def gradient_descent(b, A, initial_x, max_iters, gamma):
2.     """Gradient descent algorithm."""
3.     # Define parameters to store x and objective func. values
4.     xs = [initial_x]
5.     objectives = []
6.     x = initial_x
7.     for n_iter in range(max_iters):
8.         # *****
9.         # INSERT YOUR CODE HERE
10.        # TODO: compute gradient and objective function
11.        # *****
12.        grad, Axmb = compute_gradient(b, A, x)
13.        obj = 0.5 * np.sum(Axmb ** 2) / len(Axmb)
14.        # *****
15.        # INSERT YOUR CODE HERE
16.        # TODO: update x by a gradient descent step
17.        # *****
18.        x = x - gamma * grad
19.        # store x and objective function value
20.        xs.append(x)
21.        objectives.append(obj)
22.        print("Gradient Descent({bi}/{ti}): objective={l}".format(
23.            bi=n_iter, ti=max_iters - 1, l=obj))
24.
25.    return objectives, xs
```

Test your gradient descent function with a naive step size through gradient descent demo shown below:

```
1. # from gradient_descent import *
2. from plots import gradient_descent_visualization
```

```

3.
4. # Define the parameters of the algorithm.
5. max_iters = 50
6.
7. gamma = 0.85
8.
9. # Initialization
10. x_initial = np.zeros(A.shape[1])
11.
12. # Start gradient descent.
13. start_time = datetime.datetime.now()
14. gradient_objectives_naive, gradient_xs_naive = gradient_descent(b, A, x_initial,
    max_iters, gamma)
15. end_time = datetime.datetime.now()
16.
17. # Print result
18. exection_time = (end_time - start_time).total_seconds()
19. print("Gradient Descent: execution time={t:.3f} seconds".format(t=exection_time))

```

```

1. Gradient Descent(0/49): objective=2792.2367127591674
2. Gradient Descent(1/49): objective=77.86503142886588
3. Gradient Descent(2/49): objective=16.791668598930205
4. Gradient Descent(3/49): objective=15.417517935256669
5. Gradient Descent(4/49): objective=15.386599545324014
6. Gradient Descent(5/49): objective=15.385903881550528
7. Gradient Descent(6/49): objective=15.385888229115624
8. Gradient Descent(7/49): objective=15.385887876935838
9. Gradient Descent(8/49): objective=15.385887869011794
10. Gradient Descent(9/49): objective=15.385887868833503
11. Gradient Descent(10/49): objective=15.385887868829494
12. Gradient Descent(11/49): objective=15.385887868829403
13. Gradient Descent(12/49): objective=15.3858878688294
14. Gradient Descent(13/49): objective=15.385887868829398
15. Gradient Descent(14/49): objective=15.3858878688294
16. Gradient Descent(15/49): objective=15.3858878688294
17. Gradient Descent(16/49): objective=15.385887868829398
18. Gradient Descent(17/49): objective=15.3858878688294
19. Gradient Descent(18/49): objective=15.3858878688294
20. Gradient Descent(19/49): objective=15.3858878688294
21. Gradient Descent(20/49): objective=15.3858878688294
22. Gradient Descent(21/49): objective=15.3858878688294
23. Gradient Descent(22/49): objective=15.3858878688294
24. Gradient Descent(23/49): objective=15.3858878688294
25. Gradient Descent(24/49): objective=15.3858878688294
26. Gradient Descent(25/49): objective=15.3858878688294
27. Gradient Descent(26/49): objective=15.3858878688294

```

```
28. Gradient Descent(27/49): objective=15.3858878688294
29. Gradient Descent(28/49): objective=15.3858878688294
30. Gradient Descent(29/49): objective=15.3858878688294
31. Gradient Descent(30/49): objective=15.3858878688294
32. Gradient Descent(31/49): objective=15.3858878688294
33. Gradient Descent(32/49): objective=15.3858878688294
34. Gradient Descent(33/49): objective=15.3858878688294
35. Gradient Descent(34/49): objective=15.3858878688294
36. Gradient Descent(35/49): objective=15.3858878688294
37. Gradient Descent(36/49): objective=15.3858878688294
38. Gradient Descent(37/49): objective=15.3858878688294
39. Gradient Descent(38/49): objective=15.3858878688294
40. Gradient Descent(39/49): objective=15.3858878688294
41. Gradient Descent(40/49): objective=15.3858878688294
42. Gradient Descent(41/49): objective=15.3858878688294
43. Gradient Descent(42/49): objective=15.3858878688294
44. Gradient Descent(43/49): objective=15.3858878688294
45. Gradient Descent(44/49): objective=15.3858878688294
46. Gradient Descent(45/49): objective=15.3858878688294
47. Gradient Descent(46/49): objective=15.3858878688294
48. Gradient Descent(47/49): objective=15.3858878688294
49. Gradient Descent(48/49): objective=15.3858878688294
50. Gradient Descent(49/49): objective=15.3858878688294
51. Gradient Descent: execution time=0.002 seconds
```

Time Visualization

```
1. from ipywidgets import interact, IntSlider
2. from grid_search import *
3.
4.
5. def plot_figure(n_iter):
6.     # Generate grid data for visualization (parameters to be swept and best
       combination)
7.     grid_x0, grid_x1 = generate_w(num_intervals=10)
8.     grid_objectives = grid_search(b, A, grid_x0, grid_x1)
9.     obj_star, x0_star, x1_star = get_best_parameters(grid_x0, grid_x1,
       grid_objectives)
10.
11.     fig = gradient_descent_visualization(
12.         gradient_objectives, gradient_xs, grid_objectives, grid_x0, grid_x1,
       mean_x, std_x, height, weight, n_iter)
13.     fig.set_size_inches(20.0, 4.0)
14.     display(fig)
15.
```

```
16. interact(plot_figure, n_iter=IntSlider(min=1, max=len(gradient_xs_naive)))
```

```
1. interactive(children=(IntSlider(value=1, description='n_iter', max=51, min=1),
    Output()), _dom_classes=('widge...
2.
3.
4.
5.
6.
7. <function __main__.plot_figure(n_iter)>
```

Try doing gradient descent with a better learning rate

```
1. # Define the parameters of the algorithm.
2. max_iters = 50
3.
4. # *****
5. # INSERT YOUR CODE HERE
6. # TODO: a better learning rate using the smoothness of f
7. # *****
8. L = np.linalg.norm(A.T @ A, 2)
9. gamma = 100 / L
10.
11. # Initialization
12. x_initial = np.zeros(A.shape[1])
13.
14. # Start gradient descent.
15. start_time = datetime.datetime.now()
16. gradient_objectives, gradient_xs = gradient_descent(b, A, x_initial, max_iters,
    gamma)
17. end_time = datetime.datetime.now()
18.
19. # Print result
20. exection_time = (end_time - start_time).total_seconds()
21. print("Gradient Descent: execution time={t:.3f} seconds".format(t=exection_time))
```

```
1. Gradient Descent(0/49): objective=780.8686016504854
2. Gradient Descent(1/49): objective=773.251154488717
3. Gradient Descent(2/49): objective=766.3890713588925
4. Gradient Descent(3/49): objective=760.2039370530574
5. Gradient Descent(4/49): objective=754.6257563657889
6. Gradient Descent(5/49): objective=749.5920220920866
7. Gradient Descent(6/49): objective=745.046889499426
8. Gradient Descent(7/49): objective=740.9404446521098
9. Gradient Descent(8/49): objective=737.2280555322218
```

```
10. Gradient Descent(9/49): objective=733.8697962622064
11. Gradient Descent(10/49): objective=730.8299359180801
12. Gradient Descent(11/49): objective=728.0764844539572
13. Gradient Descent(12/49): objective=725.580789158749
14. Gradient Descent(13/49): objective=723.317175852385
15. Gradient Descent(14/49): objective=721.2626297169031
16. Gradient Descent(15/49): objective=719.3965112603233
17. Gradient Descent(16/49): objective=717.7003034395566
18. Gradient Descent(17/49): objective=716.1573864323553
19. Gradient Descent(18/49): objective=714.7528369557789
20. Gradient Descent(19/49): objective=713.47324938702
21. Gradient Descent(20/49): objective=712.3065762579041
22. Gradient Descent(21/49): objective=711.241985972315
23. Gradient Descent(22/49): objective=710.2697358408802
24. Gradient Descent(23/49): objective=709.3810587435117
25. Gradient Descent(24/49): objective=708.5680619213746
26. Gradient Descent(25/49): objective=707.8236365686215
27. Gradient Descent(26/49): objective=707.1413770434442
28. Gradient Descent(27/49): objective=706.5155086500336
29. Gradient Descent(28/49): objective=705.9408230599074
30. Gradient Descent(29/49): objective=705.4126205446067
31. Gradient Descent(30/49): objective=704.9266582834891
32. Gradient Descent(31/49): objective=704.4791040917067
33. Gradient Descent(32/49): objective=704.0664949855989
34. Gradient Descent(33/49): objective=703.6857000667632
35. Gradient Descent(34/49): objective=703.333887262911
36. Gradient Descent(35/49): objective=703.0084935140999
37. Gradient Descent(36/49): objective=702.7071980377789
38. Gradient Descent(37/49): objective=702.42789834596
39. Gradient Descent(38/49): objective=702.1686887232516
40. Gradient Descent(39/49): objective=701.927840906029
41. Gradient Descent(40/49): objective=701.703786731037
42. Gradient Descent(41/49): objective=701.495102546697
43. Gradient Descent(42/49): objective=701.3004952025822
44. Gradient Descent(43/49): objective=701.1187894523173
45. Gradient Descent(44/49): objective=700.9489166227798
46. Gradient Descent(45/49): objective=700.7899044181594
47. Gradient Descent(46/49): objective=700.6408677414381
48. Gradient Descent(47/49): objective=700.5010004283125
49. Gradient Descent(48/49): objective=700.3695677996996
50. Gradient Descent(49/49): objective=700.2458999488812
51. Gradient Descent: execution time=0.001 seconds
```

Time visualization with a better learning rate

```
1. def plot_figure(n_iter):
```

```

2.     # Generate grid data for visualization (parameters to be swept and best combination)
3.     grid_x0, grid_x1 = generate_w(num_intervals=10)
4.     grid_objectives = grid_search(b, A, grid_x0, grid_x1)
5.     obj_star, x0_star, x1_star = get_best_parameters(grid_x0, grid_x1,
        grid_objectives)
6.
7.     fig = gradient_descent_visualization(
8.         gradient_objectives, gradient_xs, grid_objectives, grid_x0, grid_x1,
        mean_x, std_x, height, weight, n_iter)
9.     fig.set_size_inches(10.0, 6.0)
10.    display(fig)
11.
12. interact(plot_figure, n_iter=IntSlider(min=1, max=len(gradient_xs)))

```

```

1. interactive(children=(IntSlider(value=1, description='n_iter', max=51, min=1),
    Output()), _dom_classes=('widget...
2.
3.
4.
5.
6.
7. <function __main__.plot_figure(n_iter)>

```

Loading more complex data

The data is taken from <https://archive.ics.uci.edu/ml/datasets/Concrete+Compressive+Strength>

```

1. data = np.loadtxt("Concrete_Data.csv", delimiter=",")
2.
3. A = data[:, :-1]
4. b = data[:, -1]
5. A, mean_A, std_A = standardize(A)

```

```

1. print('Number of samples n = ', b.shape[0])
2. print('Dimension of each sample d = ', A.shape[1])

```

```

1. Number of samples n = 1030
2. Dimension of each sample d = 8

```


Running gradient descent

Assuming bounded gradients

Assume we are moving in a bounded region $\|x\| \leq 25$ containing all iterates (and we assume $\|x - x^*\| \leq 25$ as well, for simplicity). Then by $\|\nabla f(x)\| = \frac{1}{n} \|A^T (Ax - b)\|$, one can see that f is Lipschitz over that bounded region, with Lipschitz constant $\|\nabla f(x)\| \leq \frac{1}{n} (\|A^T A\| \|x\| + \|A^T b\|)$

```
1. # *****
2. # INSERT YOUR CODE HERE
3. # TODO: Compute the bound on the gradient norm
4. # *****
5. norm_ATA = np.linalg.norm(A.T @ A, 2)
6. norm_ATb = np.linalg.norm(A.T @ b, 2)
7. n = b.shape[0]
8. grad_norm_bound = (norm_ATA * 25 + norm_ATb) / n
```

Fill in the learning rate assuming bounded gradients

```
1. max_iters = 50
2.
3. # *****
4. # INSERT YOUR CODE HERE
5. # TODO: Compute learning rate based on bounded gradient
6. # *****
7.
8. gamma = 10 / grad_norm_bound
9. print("grad_norm_bound:", grad_norm_bound)
10. print("gamma:", gamma)
11.
12. # Initialization
13. x_initial = np.zeros(A.shape[1])
14.
15. # Start gradient descent.
16. start_time = datetime.datetime.now()
17. bd_gradient_objectives, bd_gradient_xs = gradient_descent(b, A, x_initial,
    max_iters, gamma)
18. end_time = datetime.datetime.now()
19.
20.
21. # Print result
22. execution_time = (end_time - start_time).total_seconds()
23. print("Gradient Descent: execution time={t:.3f} seconds".format(t=execution_time))
```

```

24.
25. # Averaging the iterates as is the case for bounded gradients case
26. bd_gradient_objectives_averaged = []
27. for i in range(len(bd_gradient_xs)):
28.     if i > 0:
29.         bd_gradient_xs[i] = (i * bd_gradient_xs[i-1] + bd_gradient_xs[i])/(i + 1)
30.     grad, err = compute_gradient(b, A, bd_gradient_xs[i])
31.     obj = calculate_objective(err)
32.     bd_gradient_objectives_averaged.append(obj)

```

```

1. grad_norm_bound: 70.55157320349318
2. gamma: 0.14174028368094385
3. Gradient Descent(0/49): objective=780.8686016504854
4. Gradient Descent(1/49): objective=757.0572465196063
5. Gradient Descent(2/49): objective=740.5812543624813
6. Gradient Descent(3/49): objective=729.084584479613
7. Gradient Descent(4/49): objective=720.9907413148369
8. Gradient Descent(5/49): objective=715.2382349171139
9. Gradient Descent(6/49): objective=711.1080185774268
10. Gradient Descent(7/49): objective=708.1100789289962
11. Gradient Descent(8/49): objective=705.9084315994309
12. Gradient Descent(9/49): objective=704.2712418678465
13. Gradient Descent(10/49): objective=703.0374846206412
14. Gradient Descent(11/49): objective=702.0945497450593
15. Gradient Descent(12/49): objective=701.3631255793681
16. Gradient Descent(13/49): objective=700.7869435574245
17. Gradient Descent(14/49): objective=700.3257840958671
18. Gradient Descent(15/49): objective=699.9506801797878
19. Gradient Descent(16/49): objective=699.6406088594872
20. Gradient Descent(17/49): objective=699.3801950738394
21. Gradient Descent(18/49): objective=699.1581078299688
22. Gradient Descent(19/49): objective=698.9659325325084
23. Gradient Descent(20/49): objective=698.7973726954689
24. Gradient Descent(21/49): objective=698.6476809154398
25. Gradient Descent(22/49): objective=698.5132504450324
26. Gradient Descent(23/49): objective=698.3913200155959
27. Gradient Descent(24/49): objective=698.2797590603523
28. Gradient Descent(25/49): objective=698.1769104076669
29. Gradient Descent(26/49): objective=698.0814743343121
30. Gradient Descent(27/49): objective=697.992422585014
31. Gradient Descent(28/49): objective=697.9089342458274
32. Gradient Descent(29/49): objective=697.8303476560615
33. Gradient Descent(30/49): objective=697.7561241621568
34. Gradient Descent(31/49): objective=697.6858206650937
35. Gradient Descent(32/49): objective=697.6190687328028
36. Gradient Descent(33/49): objective=697.5555586384422

```

```

37. Gradient Descent(34/49): objective=697.495027111916
38. Gradient Descent(35/49): objective=697.4372479026222
39. Gradient Descent(36/49): objective=697.3820244790351
40. Gradient Descent(37/49): objective=697.329184358537
41. Gradient Descent(38/49): objective=697.2785746852895
42. Gradient Descent(39/49): objective=697.2300587666316
43. Gradient Descent(40/49): objective=697.183513347901
44. Gradient Descent(41/49): objective=697.1388264577861
45. Gradient Descent(42/49): objective=697.0958956957468
46. Gradient Descent(43/49): objective=697.0546268629419
47. Gradient Descent(44/49): objective=697.0149328608449
48. Gradient Descent(45/49): objective=696.9767327990826
49. Gradient Descent(46/49): objective=696.9399512673167
50. Gradient Descent(47/49): objective=696.9045177361759
51. Gradient Descent(48/49): objective=696.8703660600796
52. Gradient Descent(49/49): objective=696.8374340608425
53. Gradient Descent: execution time=0.002 seconds

```

Gradient descent using smoothness

Fill in the learning rate using smoothness of the function

```

1. max_iters = 50
2.
3.
4. # *****
5. # INSERT YOUR CODE HERE
6. # TODO: a better learning rate using the smoothness of f
7. # *****
8. L = np.linalg.norm(A.T @ A, 2)
9. gamma = 100 / L
10.
11. # Initialization
12. x_initial = np.zeros(A.shape[1])
13.
14. # Start gradient descent.
15. start_time = datetime.datetime.now()
16. gradient_objectives, gradient_xs = gradient_descent(b, A, x_initial, max_iters,
    gamma)
17. end_time = datetime.datetime.now()
18.
19. # Print result
20. exection_time = (end_time - start_time).total_seconds()
21. print("Gradient Descent: execution time={t:.3f} seconds".format(t=exection_time))

```

1. **Gradient Descent**(0/49): objective=780.8686016504854
2. **Gradient Descent**(1/49): objective=773.251154488717
3. **Gradient Descent**(2/49): objective=766.3890713588925
4. **Gradient Descent**(3/49): objective=760.2039370530574
5. **Gradient Descent**(4/49): objective=754.6257563657889
6. **Gradient Descent**(5/49): objective=749.5920220920866
7. **Gradient Descent**(6/49): objective=745.046889499426
8. **Gradient Descent**(7/49): objective=740.9404446521098
9. **Gradient Descent**(8/49): objective=737.2280555322218
10. **Gradient Descent**(9/49): objective=733.8697962622064
11. **Gradient Descent**(10/49): objective=730.8299359180801
12. **Gradient Descent**(11/49): objective=728.0764844539572
13. **Gradient Descent**(12/49): objective=725.580789158749
14. **Gradient Descent**(13/49): objective=723.317175852385
15. **Gradient Descent**(14/49): objective=721.2626297169031
16. **Gradient Descent**(15/49): objective=719.3965112603233
17. **Gradient Descent**(16/49): objective=717.7003034395566
18. **Gradient Descent**(17/49): objective=716.1573864323553
19. **Gradient Descent**(18/49): objective=714.7528369557789
20. **Gradient Descent**(19/49): objective=713.47324938702
21. **Gradient Descent**(20/49): objective=712.3065762579041
22. **Gradient Descent**(21/49): objective=711.241985972315
23. **Gradient Descent**(22/49): objective=710.2697358408802
24. **Gradient Descent**(23/49): objective=709.3810587435117
25. **Gradient Descent**(24/49): objective=708.5680619213746
26. **Gradient Descent**(25/49): objective=707.8236365686215
27. **Gradient Descent**(26/49): objective=707.1413770434442
28. **Gradient Descent**(27/49): objective=706.5155086500336
29. **Gradient Descent**(28/49): objective=705.9408230599074
30. **Gradient Descent**(29/49): objective=705.4126205446067
31. **Gradient Descent**(30/49): objective=704.9266582834891
32. **Gradient Descent**(31/49): objective=704.4791040917067
33. **Gradient Descent**(32/49): objective=704.0664949855989
34. **Gradient Descent**(33/49): objective=703.6857000667632
35. **Gradient Descent**(34/49): objective=703.333887262911
36. **Gradient Descent**(35/49): objective=703.0084935140999
37. **Gradient Descent**(36/49): objective=702.7071980377789
38. **Gradient Descent**(37/49): objective=702.42789834596
39. **Gradient Descent**(38/49): objective=702.1686887232516
40. **Gradient Descent**(39/49): objective=701.927840906029
41. **Gradient Descent**(40/49): objective=701.703786731037
42. **Gradient Descent**(41/49): objective=701.495102546697
43. **Gradient Descent**(42/49): objective=701.3004952025822
44. **Gradient Descent**(43/49): objective=701.1187894523173
45. **Gradient Descent**(44/49): objective=700.9489166227798
46. **Gradient Descent**(45/49): objective=700.7899044181594

```
47. Gradient Descent(46/49): objective=700.6408677414381
48. Gradient Descent(47/49): objective=700.5010004283125
49. Gradient Descent(48/49): objective=700.3695677996996
50. Gradient Descent(49/49): objective=700.2458999488812
51. Gradient Descent: execution time=0.001 seconds
```

Plotting the Evolution of the Objective Function

```
1. plt.figure(figsize=(8, 8))
2. plt.xlabel('Number of steps')
3. plt.ylabel('Objective Function')
4. #plt.yscale("log")
5. plt.plot(range(len(gradient_objectives)), gradient_objectives, 'r',
            label='gradient descent with 1/L stepsize')
6. plt.plot(range(len(bd_gradient_objectives)), bd_gradient_objectives, 'b',
            label='gradient descent assuming bounded gradients')
7. plt.plot(range(len(bd_gradient_objectives_averaged)),
            bd_gradient_objectives_averaged, 'g', label='gradient descent assuming bounded
            gradients with averaged iterates')
8. plt.legend(loc='upper right')
9. plt.show()
```

